

Spring 2024:CS5720 Neural Networks & Deep Learning

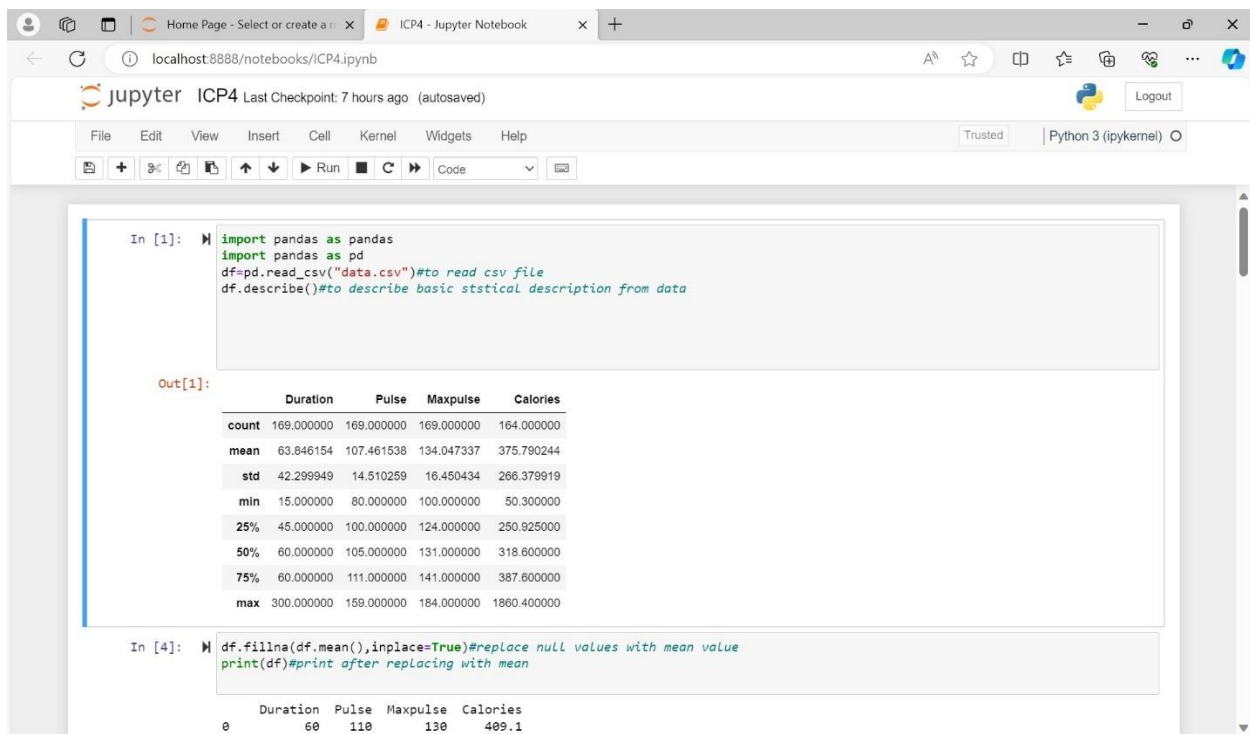
ICP-4 Assignment-4

Name: Sai Sushma Sri Bireddy Student ID:700747557

GitHub Link: <https://github.com/SaiSushmaSriBireddy/Assignment4>

Video Link: <https://drive.google.com/file/d/11OtYiAjpB-2fSCnroVmZvlvsKKvCrZvv/view?usp=sharing>

- Read the provided CSV file 'data.csv'.
- <https://drive.google.com/drive/folders/1h8C3mLsso-R-slOLsvoYwPLzy2fJ4IOF?usp=sharing>
- Show the basic statistical description about the data.



The screenshot shows a Jupyter Notebook interface with the following content:

```
In [1]: import pandas as pandas
import pandas as pd
df=pd.read_csv("data.csv")#to read csv file
df.describe()#to describe basic ststical description from data
```

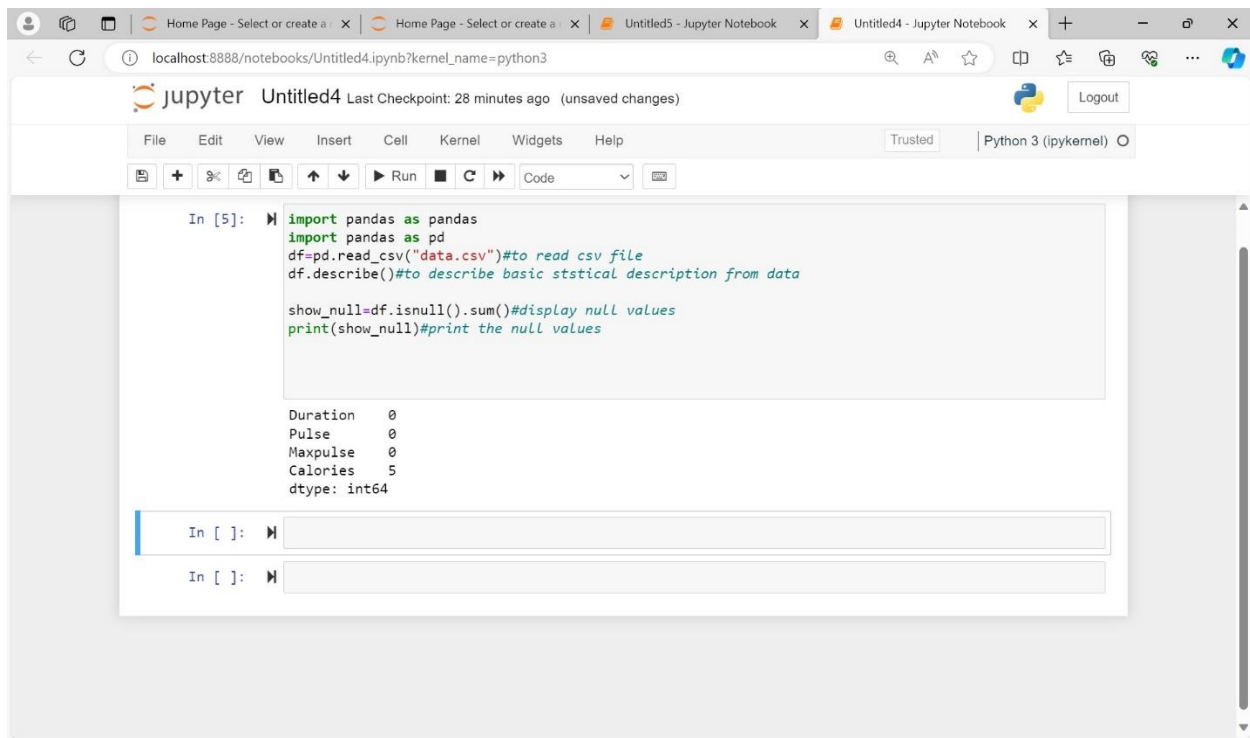
Out[1]:

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

```
In [4]: df.fillna(df.mean(),inplace=True)#replace null values with mean value
print(df)#print after replacing with mean
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1

- Check if the data has null values. i. Replace the null values with the mean



localhost8888/notebooks/Untitled4.ipynb?kernel_name=python3

jupyter Untitled4 Last Checkpoint: 28 minutes ago (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

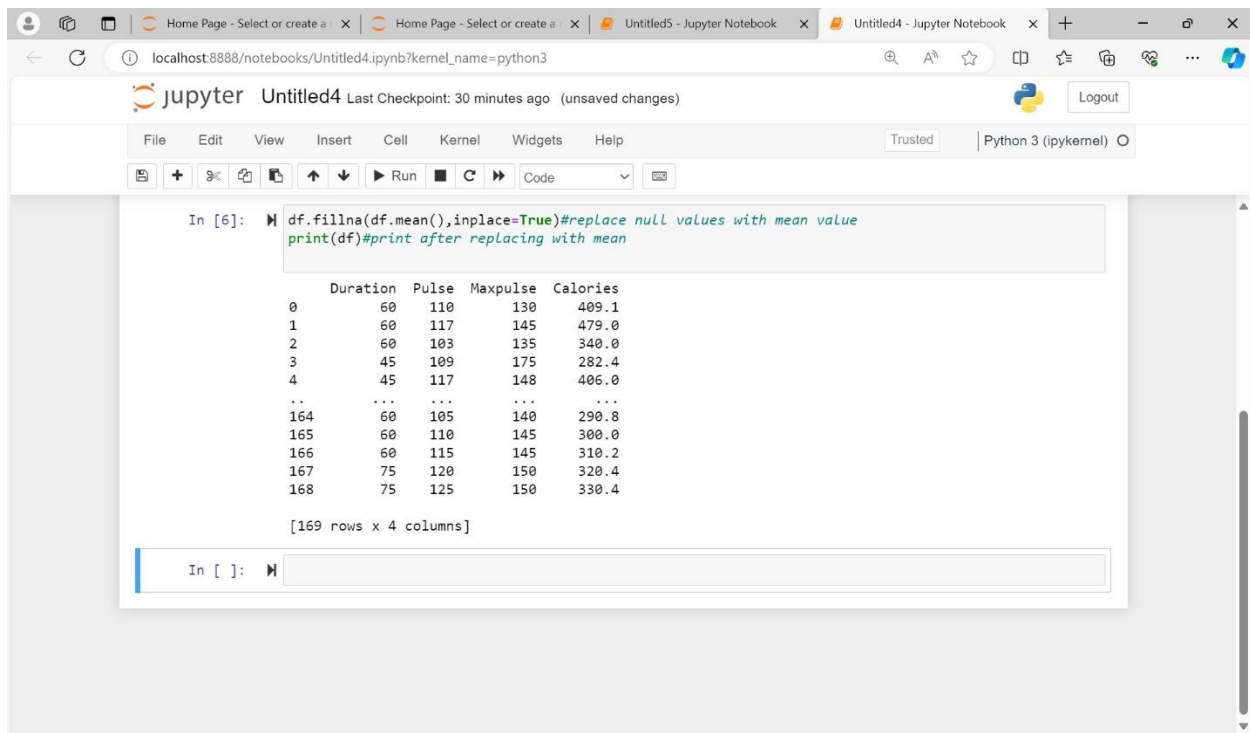
```
In [5]: import pandas as pandas
import pandas as pd
df=pd.read_csv("data.csv")#to read csv file
df.describe()#to describe basic ststical description from data

show_null=df.isnull().sum()#display null values
print(show_null)#print the null values
```

```
Duration    0
Pulse       0
Maxpulse    0
Calories    5
dtype: int64
```

In []:

In []:



localhost8888/notebooks/Untitled4.ipynb?kernel_name=python3

jupyter Untitled4 Last Checkpoint: 30 minutes ago (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [6]: df.fillna(df.mean(),inplace=True)#replace null values with mean value
print(df)#print after replacing with mean
```

```
   Duration  Pulse  Maxpulse  Calories
0         60    110      130    409.1
1         60    117      145    479.0
2         60    103      135    340.0
3         45    109      175    282.4
4         45    117      148    406.0
...      ...    ...      ...      ...
164        60    105      140    290.8
165        60    110      145    300.0
166        60    115      145    310.2
167        75    120      150    320.4
168        75    125      150    330.4

[169 rows x 4 columns]
```

In []:

e. Select at least two columns and aggregate the data using: min, max, count, mean.

The screenshot shows a Jupyter Notebook interface with the following code and output:

```
df=df[["Duration","Pulse","Maxpulse","Calories"]]
aggregate={"Duration":["max","min","count","mean"],
           "Pulse":["max","min","count","mean"],
           "Maxpulse":["max","min","count","mean"],
           "Calories":["max","min","count","mean"]}# to find max,min,count,mean of all the columns
aggregate_df=df.agg(aggregate)#function to aggregate
print(aggregate_df)#print the aggregate
```

	Duration	Pulse	Maxpulse	Calories
max	300.000000	159.000000	184.000000	1860.400000
min	15.000000	80.000000	100.000000	50.300000
count	169.000000	169.000000	169.000000	169.000000
mean	63.846154	107.461538	134.047337	375.790244

f. Filter the data frame to select the rows with calories values between 500 and 1000.

The screenshot shows a Jupyter Notebook interface with the following code and output:

```
calories_in_range=(df["Calories"]>=500) & (df["Calories"]<=1000)#defining range of values to be displayed
filters_result=df[calories_in_range]#adding the defined range to new variable
print(filters_result)#printing the new result
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
83	120	100	130	500.0
90	180	101	127	600.1
99	90	93	124	604.1
101	90	90	110	500.0
102	90	90	100	500.0
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

g. Filter the data frame to select the rows with calories values > 500 and pulse < 100.

A Jupyter Notebook interface with the title "Untitled4" and a last checkpoint of 32 minutes ago. The notebook contains a code cell with the following Python code:

```
calories_pulse_filter=(df["Calories"]>500)&(df["Pulse"]<100)#defining range
filters_result=df[calories_pulse_filter]#adding the result to new variable
print(filters_result)#printing the result
```

The output of the code is a DataFrame with 5 columns: Duration, Pulse, Maxpulse, and Calories. The data is as follows:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

h. Create a new “df_modified” dataframe that contains all the columns from df except for “Maxpulse”.

A Jupyter Notebook interface with the title "Untitled4" and a last checkpoint of 33 minutes ago (autosaved). The notebook contains a code cell with the following Python code:

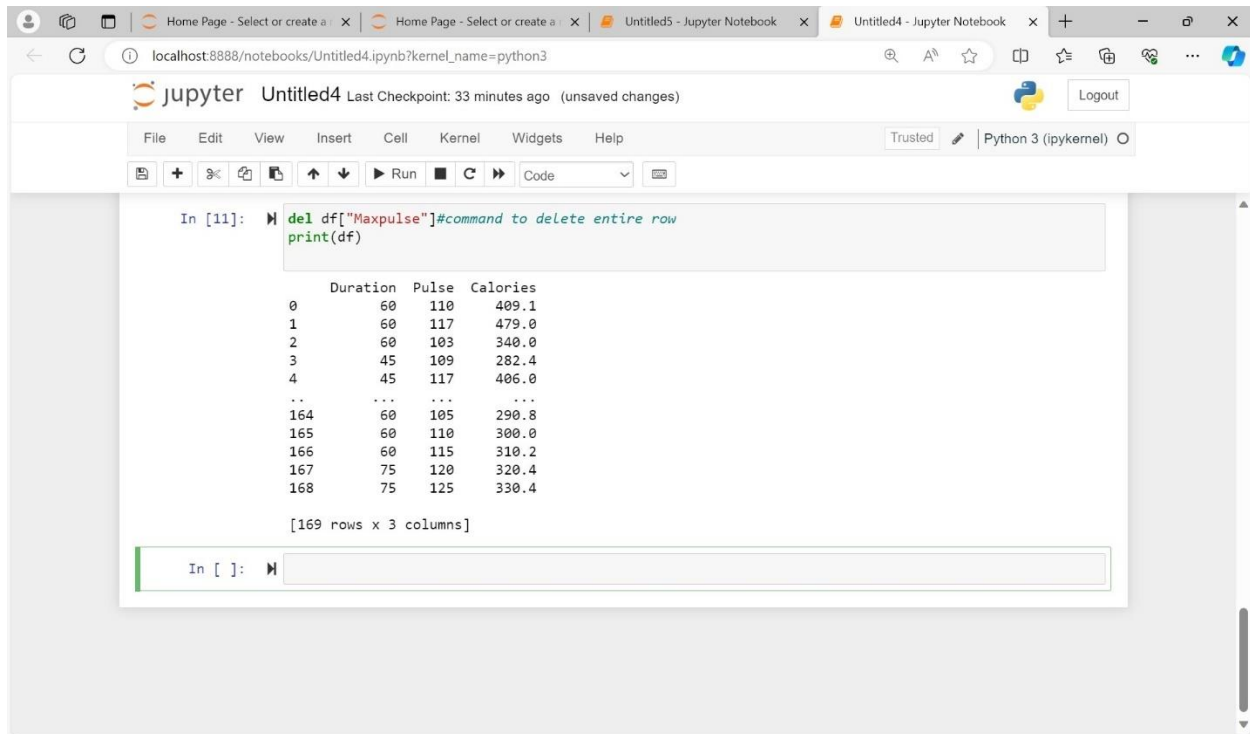
```
df_modified=df.drop(columns=["Maxpulse"])#displaying every column except Maxpulse
print(df_modified)#printing the result
```

The output of the code is a DataFrame with 3 columns: Duration, Pulse, and Calories. The data is as follows:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0
...
164	60	105	290.8
165	60	110	300.0
166	60	115	310.2
167	75	120	320.4
168	75	125	330.4

The output also includes the text "[169 rows x 3 columns]" at the bottom.

I. Delete the “Maxpulse” column from the main df data frame



A screenshot of a Jupyter Notebook interface. The browser tabs at the top show 'Home Page - Select or create a...', 'Untitled5 - Jupyter Notebook', and 'Untitled4 - Jupyter Notebook'. The address bar shows 'localhost:8888/notebooks/Untitled4.ipynb?kernel_name=python3'. The notebook title is 'jupyter Untitled4' with a subtitle 'Last Checkpoint: 33 minutes ago (unsaved changes)'. The menu bar includes 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. The toolbar has icons for file operations, a 'Run' button, and a dropdown menu set to 'Code'. The code cell contains the following Python code:

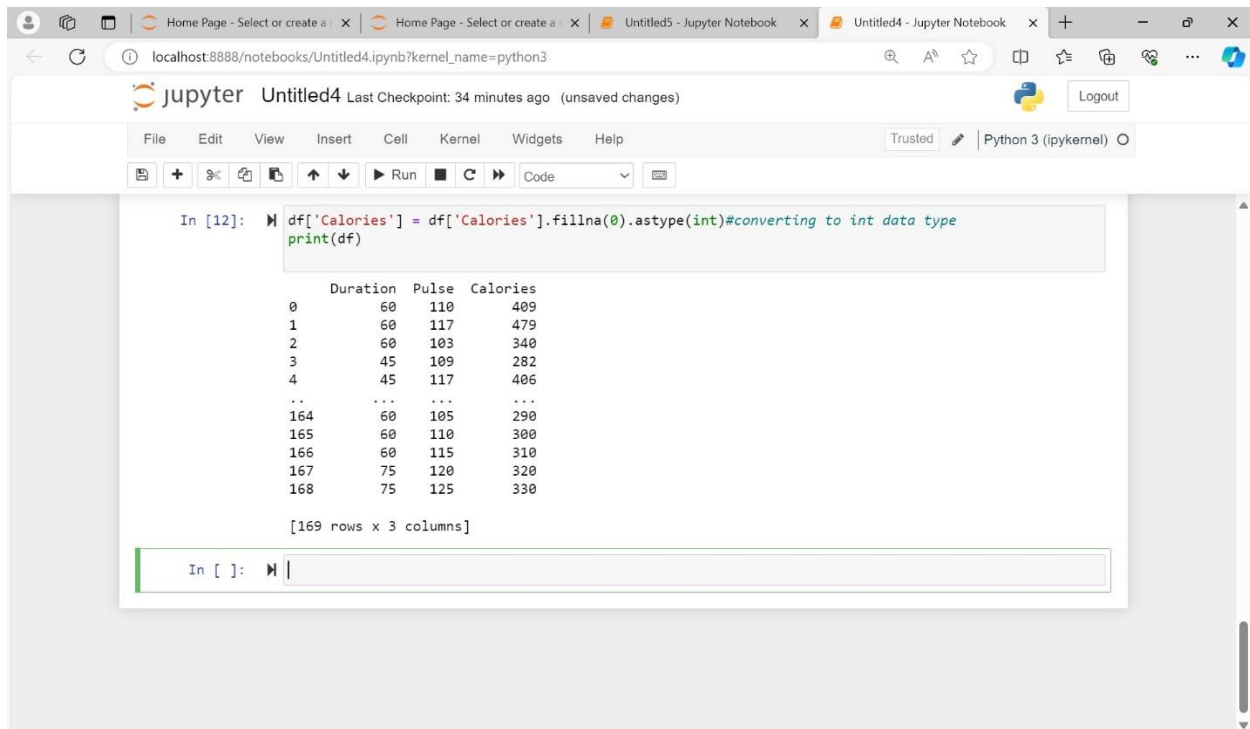
```
In [11]: del df["Maxpulse"]#command to delete entire row
print(df)
```

The output of the code is a DataFrame with 169 rows and 3 columns: 'Duration', 'Pulse', and 'Calories'. The data is as follows:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0
...
164	60	105	290.8
165	60	110	300.0
166	60	115	310.2
167	75	120	320.4
168	75	125	330.4

Below the output, it says '[169 rows x 3 columns]'. The input prompt 'In []:' is visible at the bottom of the cell.

j. Convert the datatype of Calories column to int datatype.



A screenshot of a Jupyter Notebook interface, similar to the one above. The browser tabs and address bar are the same. The notebook title is 'jupyter Untitled4' with a subtitle 'Last Checkpoint: 34 minutes ago (unsaved changes)'. The code cell contains the following Python code:

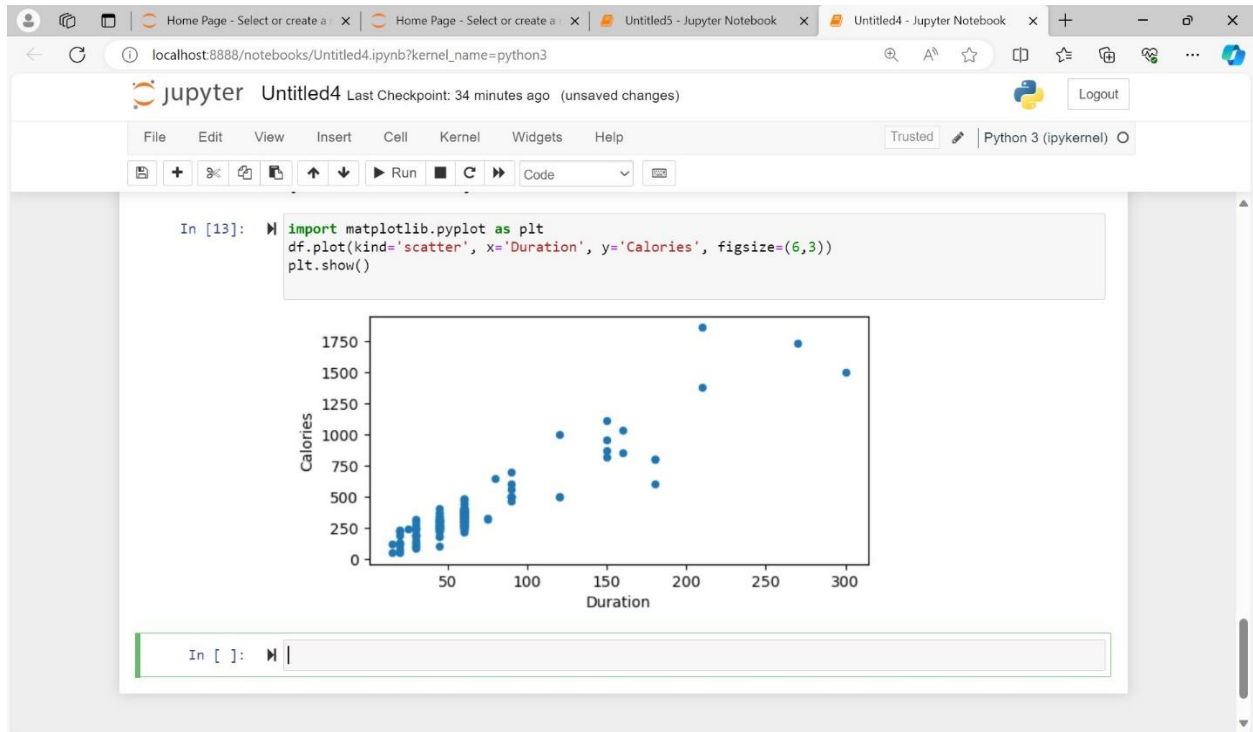
```
In [12]: df['Calories'] = df['Calories'].fillna(0).astype(int)#converting to int data type
print(df)
```

The output of the code is a DataFrame with 169 rows and 3 columns: 'Duration', 'Pulse', and 'Calories'. The data is as follows:

	Duration	Pulse	Calories
0	60	110	409
1	60	117	479
2	60	103	340
3	45	109	282
4	45	117	406
...
164	60	105	290
165	60	110	300
166	60	115	310
167	75	120	320
168	75	125	330

Below the output, it says '[169 rows x 3 columns]'. The input prompt 'In []:' is visible at the bottom of the cell.

k. Using pandas create a scatter plot for the two columns (Duration and Calories).



2. Linear Regression

- Import the given "Salary_Data.csv"
- Split the data in train test partitions, such that 1/3 of the data is reserved as test subset.
- Train and predict the model.
- Calculate the mean squared error
- Visualize both train and test data using scatter plot.

```
In [10]: import pandas as pd
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
salariesData = pd.read_csv('Salary_Data (2).csv') #importing data from the CSV file
df.describe()
#splitting the data in to training and testing
X = salariesData.iloc[:, :-1].values
Y = salariesData.iloc[:, 1].values
#splitting 1/3 of the data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 1/3, random_state = 0)
# Fitting Simple Linear Regression to the training set
reg = LinearRegression()
reg.fit(X_train, Y_train)
# Predicting the Test set result
pred = reg.predict(X_test)
# Calculating the Mean_squared_error
mse = mean_squared_error(Y_test, pred)
#Visualising the Training set results and Test set results
plt.scatter(X_train, Y_train, color = 'blue')
plt.scatter(X_test, Y_test, color = 'red')
plt.title('Salary Data')
plt.xlabel('Experience (Years)')
plt.ylabel('Salary')
plt.show()
```

Output:

